

# Stat 435 lecture notes 9: supplementary

## Penalized logistic regression

```
load("sim_gen0.rda")
load("sim_trait.rda")
# take a smaller par
sim_gen0_a = sim_gen0[1:500, 1:100]
sim_trait_a = sim_trait[1:100]
sim_trait_a1 = as.numeric(abs(sim_trait_a) > 0.5)

set.seed(123)

library(glmnet)

lga = cv.glmnet(t(sim_gen0_a), sim_trait_a1, family = "binomial",
  alpha = 1, type.measure = "class")

Ma = glmnet(t(sim_gen0_a), sim_trait_a1, alpha = 1, family = "binomial",
  lambda = lga$lambda.min)
# Display non-zero regression coefficients
coef(Ma)[coef(Ma) != 0]

## <sparse>[ <logic> ] : .M.sub.i.logical() maybe inefficient

## [1] 0.906593001 -0.050484903 -0.142419280 0.198313448 -0.055539426
## [6] -0.548750948 0.070020997 0.040914023 0.855089233 0.292877286
## [11] 0.565699508 -0.097575022 0.048652471 0.077754855 -0.142844461
## [16] -0.463245195 -0.236106215 -0.273159753 -0.569664748 -0.067262219
## [21] -0.332920025 -0.527457719 -1.008362016 0.053478316 -0.098952564
## [26] 0.311832251 -0.229191419 -0.611895254 0.048090837 0.506896889
## [31] 0.084019436 0.008013738 0.263433125 -0.114768149 0.685211115
## [36] 0.559790810 -0.754621693 0.018782008 0.460460521 -0.118559394
## [41] 0.508248393 -0.043992653

## obtain p-values
library(hdi)

## Warning: package 'hdi' was built under R version 3.5.3

## Loading required package: scalreg

## Warning: package 'scalreg' was built under R version 3.5.2

## Loading required package: lars

## Warning: package 'lars' was built under R version 3.5.2

## Loaded lars 1.2
```

```

prj_est = lasso.proj(t(sim_geno_a), sim_trait_a1, family = "binomial")

## Nodewise regressions will be computed as no argument Z was provided.
## You can store Z to avoid the majority of the computation next time around.
## Z only depends on the design matrix x.

## Warning in warning.sigma.message(): Overriding the error variance estimate
## with your own value. The initial estimate implies an error variance
## estimate and if they don't correspond the testing might not be correct
## anymore.

prj_est$pval[1:10]

## [1] 0.2981953 0.7763949 0.5395813 0.8914452 0.5055726 0.5910186 0.9138428
## [8] 0.8536747 0.9884205 0.4794335

### Ridge regression

lgb = cv.glmnet(t(sim_geno_a), sim_trait_a1, family = "binomial",
  alpha = 0, type.measure = "class")

Mb = glmnet(t(sim_geno_a), sim_trait_a1, alpha = 0, family = "binomial",
  lambda = lga$lambda.min)

## obtain p-values
prj_estb = ridge.proj(t(sim_geno_a), sim_trait_a1, family = "binomial")

## Warning in warning.sigma.message(): Overriding the error variance estimate
## with your own value. The initial estimate implies an error variance
## estimate and if they don't correspond the testing might not be correct
## anymore.

prj_estb$pval[1:10]

## [1] 0.9594145 1.0000000 1.0000000 1.0000000 0.9233757 0.9354331 1.0000000
## [8] 1.0000000 1.0000000 0.9020482

```

## Multinomial logistic regression

Please go to webpage: <https://stats.idre.ucla.edu/r/dae/multinomial-logistic-regression/>

The data set contains variables on 200 students. The outcome variable is `prog`, program type. The predictor variables are social economic status, `ses`, a three-level categorical variable and writing score, `write`, a continuous variable.

```

require(foreign)
require(nnet)
require(ggplot2)

```

```
require(reshape2)

ml <- read.dta("hsbdemo.dta")

with(ml, table(ses, prog))

##           prog
## ses      general academic vocation
##  low           16          19          12
## middle          20          44          31
##  high           9          42           7

# set baseline level for 'prog'
ml$prog2 <- relevel(ml$prog, ref = "academic")
test <- multinom(prog2 ~ ses + write, data = ml)

## # weights:  15 (8 variable)
## initial  value 219.722458
## iter  10 value 179.982880
## final   value 179.981726
## converged

# 'ses = low' is the baseline and absorbed into the intercept
summary(test)

## Call:
## multinom(formula = prog2 ~ ses + write, data = ml)
##
## Coefficients:
##           (Intercept) sesmiddle  seshigh      write
## general      2.852198 -0.5332810 -1.1628226 -0.0579287
## vocation     5.218260  0.2913859 -0.9826649 -0.1136037
##
## Std. Errors:
##           (Intercept) sesmiddle  seshigh      write
## general      1.166441 0.4437323 0.5142196 0.02141097
## vocation     1.163552 0.4763739 0.5955665 0.02221996
##
## Residual Deviance: 359.9635
## AIC: 375.9635

# calculate z-test value
z <- summary(test)$coefficients/summary(test)$standard.errors
z

##           (Intercept) sesmiddle  seshigh      write
## general      2.445214 -1.2018081 -2.261334 -2.705562
## vocation     4.484769  0.6116747 -1.649967 -5.112689
```

```

# calculate two-sided p-values using z-tests
p <- (1 - pnorm(abs(z), 0, 1)) * 2
p

##           (Intercept) sesmiddle   seshigh      write
## general  0.0144766100 0.2294379 0.02373856 6.818902e-03
## vocation 0.0000072993 0.5407530 0.09894976 3.176045e-07

# use fitted model to predict probabilities
head(pp <- fitted(test))

##   academic   general   vocation
## 1 0.1482764 0.3382454 0.5134781
## 2 0.1202017 0.1806283 0.6991700
## 3 0.4186747 0.2368082 0.3445171
## 4 0.1726885 0.3508384 0.4764731
## 5 0.1001231 0.1689374 0.7309395
## 6 0.3533566 0.2377976 0.4088458

```

## Poisson regression

The following contents are adopted from: <https://stats.idre.ucla.edu/r/dae/poisson-regression/>. In this example, `num_awards` is the outcome variable and indicates the number of awards earned by students at a high school in a year, `math` is a continuous predictor variable and represents students' scores on their math final exam, and `prog` is a categorical predictor variable with three levels indicating the type of program in which the students were enrolled. It is coded as 1 = "General", 2 = "Academic" and 3 = "Vocational".

```

require(ggplot2)
require(sandwich)

## Loading required package: sandwich

## Warning in library(package, lib.loc = lib.loc, character.only = TRUE,
## logical.return = TRUE, : there is no package called 'sandwich'

require(msm)

## Loading required package: msm


## Warning in library(package, lib.loc = lib.loc, character.only = TRUE,
## logical.return = TRUE, : there is no package called 'msm'

p <- read.csv("poisson_sim.csv")
p <- within(p, {
  prog <- factor(prog, levels = 1:3, labels = c("General",
    "Academic", "Vocational"))
  id <- factor(id)
})
summary(p)

```

```
##          id      num_awards      prog      math
## 1      : 1  Min.    :0.00  General   : 45  Min.    :33.00
## 2      : 1  1st Qu.:0.00  Academic  :105  1st Qu.:45.00
## 3      : 1  Median :0.00  Vocational: 50  Median :52.00
## 4      : 1  Mean   :0.63                Mean   :52.65
## 5      : 1  3rd Qu.:1.00                3rd Qu.:59.00
## 6      : 1  Max.    :6.00                Max.    :75.00
## (Other):194
```

```
ggplot(p, aes(num_awards, fill = prog)) + geom_histogram(binwidth = 0.5,
  position = "dodge")
```



```
summary(m1 <- glm(num_awards ~ prog + math, family = "poisson",
  data = p))
```

```
##
## Call:
## glm(formula = num_awards ~ prog + math, family = "poisson", data = p)
##
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -2.2043  -0.8436  -0.5106   0.2558   2.6796
##
## Coefficients:
```

```
##               Estimate Std. Error z value Pr(>|z|)
## (Intercept)   -5.24712    0.65845  -7.969 1.60e-15 ***
## progAcademic    1.08386    0.35825   3.025 0.00248 **
## progVocational  0.36981    0.44107   0.838 0.40179
## math           0.07015    0.01060   6.619 3.63e-11 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for poisson family taken to be 1)
##
##      Null deviance: 287.67  on 199  degrees of freedom
## Residual deviance: 189.45  on 196  degrees of freedom
## AIC: 373.5
##
## Number of Fisher Scoring iterations: 6

## calculate and store predicted values
p$phat <- predict(m1, type = "response")

## order by program and then by math
p <- p[with(p, order(prog, math)), ]

## create the plot
ggplot(p, aes(x = math, y = phat, colour = prog)) + geom_point(aes(y = num_awards),
  alpha = 0.5, position = position_jitter(h = 0.2)) + geom_line(size = 1) +
  labs(x = "Math Score", y = "Expected number of awards")
```

