

Stat 435 Lecture Notes 3

Xiongzhi Chen
Washington State University

Contents


	2
Linear regression with a qualitative predictor	2
Motivation	2
Motivation	2
Model 1: 2 levels	2
Model 1: 2 levels	2
Fitting the Model 1	3
Testing the Model 1	3
Model 2: 3 levels	3
Model 2: 3 levels	4
Fitting the Model 2	4
Testing the Model 2	4
Testing the Model 2	5
Diagnostics	5
Multiple linear regression	6
Motivation	6
Model	6
Fitting model	6
Testing on all coefficients	6
Testing on all coefficients	7
Testing on some coefficients	7
Testing on some coefficients	7
Testing on model fit	7
Interaction terms	8
Interaction terms: I	8
Interaction terms: I	8
Interaction terms: I	8
Interaction terms: I	9
Interaction terms: II	9
Interaction terms: II	10
Interaction terms: II	10
Diagnostics	11
Diagnostics	11
Collinearity	11
Collinearity	11
Collinearity	12
Collinearity	12
Non-linear models	13
Non-linear relationship	13
License and session Information	13

Linear regression with a qualitative predictor

Motivation

- How is **Balance** of a credit card related to a user's **Gender**?
- How is **Balance** of a credit card related to a user's **Ethnicity**?

Motivation



Model 1: 2 levels

- Coding: **Gender** has 2 levels, **Male** and **Female**
- *dummy variable*: $x_i = 0$ if i th person is **Female** and $x_i = 1$ if i th person is **Male**
- Model: $y_i = \beta_0 + \beta_1 x_i + \varepsilon_i$, which induces 2 submodels:
- $y_i = \beta_0 + \beta_1 + \varepsilon_i$ if i th person is **Male**
- $y_i = \beta_0 + \varepsilon_i$ if i th person is **Female**

Note: dummy variable follows coding by R, for which the first level **Female** is the baseline

Model 1: 2 levels

Model: $y_i = \beta_0 + \beta_1 x_i + \varepsilon_i$

- $x_i = 0$ if i th person is **Female**, and $x_i = 1$ if i th person is **Male**
- β_0 : average **Balance** for females
- β_1 : average difference in balance between *males and females*

Remark: coding of a dummy variable is arbitrary and should be easily interpretable

Fitting the Model 1

Call:

```
lm(formula = Balance ~ Gender, data = creditData)
```

Coefficients:

```
(Intercept)  GenderMale
      529.54      -19.73
```

- Females have an average balance of \$529.54; Female baseline
- Males have an average balance of \$(529.54-19.73)= \$509.80

Note: in R, by default the first level **Female** is the baseline

Testing the Model 1

Call:

```
lm(formula = Balance ~ Gender, data = creditData)
```

Residuals:

```
      Min       1Q   Median       3Q      Max
-529.54 -455.35  -60.17   334.71 1489.20
```

Coefficients:

```
              Estimate Std. Error t value Pr(>|t|)
(Intercept)    529.54      31.99   16.554  <2e-16 ***
GenderMale     -19.73      46.05    -0.429    0.669
---

```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

Residual standard error: 460.2 on 398 degrees of freedom

Multiple R-squared: 0.0004611, Adjusted R-squared: -0.00205

F-statistic: 0.1836 on 1 and 398 DF, p-value: 0.6685

- If model assumptions are met, **Gender** is not significant on affecting average balance at type I error level 0.05 based on F-statistic (or the p-value of **GenderMale**)

Model 2: 3 levels

Ethnicity has 3 levels **African American** (1st level and baseline in R), **Asian**, and **Caucasian**. 2 dummy variables are needed:

- $x_{i1} = 0$ if i th person is *not Asian*, and $x_{i1} = 1$ if i th person is **Asian**
- $x_{i2} = 0$ if i th person is *not Caucasian*, and $x_{i2} = 1$ if i th person is **Caucasian**

- Model:

$$y_i = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \varepsilon_i$$

Model 2: 3 levels

Model:

$$y_i = \beta_0 + \beta_1 x_{i1} + \beta_2 x_{i2} + \varepsilon_i$$

- Codings on previous slide
- β_0 : average balance for African American
- β_1 : average difference in balance between Asian and African American
- β_2 : average difference in balance between Caucasian and African American

Fitting the Model 2

Call:

```
lm(formula = Balance ~ Ethnicity, data = creditData)
```

Coefficients:

(Intercept)	EthnicityAsian	EthnicityCaucasian
531.00	-18.69	-12.50

- African Americans have an average balance of \$531
- Asians have an average balance of \$(531-18.69)= \$512.31
- Caucasians have an average balance of \$(531-12.50)= \$518.5

Testing the Model 2

Call:

```
lm(formula = Balance ~ Ethnicity, data = creditData)
```

Residuals:

Min	1Q	Median	3Q	Max
-531.00	-457.08	-63.25	339.25	1480.50

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	531.00	46.32	11.464	<2e-16 ***
EthnicityAsian	-18.69	65.02	-0.287	0.774
EthnicityCaucasian	-12.50	56.68	-0.221	0.826

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 460.9 on 397 degrees of freedom

Multiple R-squared: 0.0002188, Adjusted R-squared: -0.004818

F-statistic: 0.04344 on 2 and 397 DF, p-value: 0.9575

If model assumptions are met, at type I error level 0.05, Ethnicity does not significantly affect average balance based on the F-statistic

Testing the Model 2

A tibble: 3 x 5

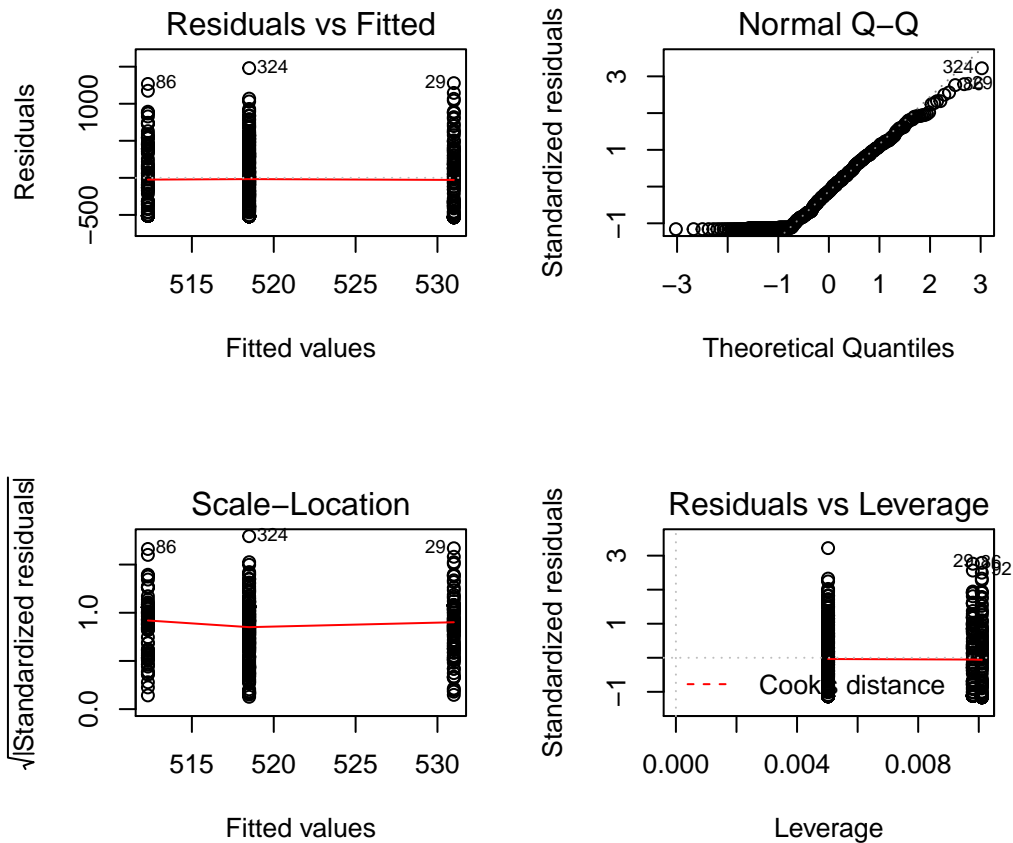
term	estimate	std.error	statistic	p.value
<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1 (Intercept)	531.	46.3	11.5	1.77e-26
2 EthnicityAsian	-18.7	65.0	-0.287	7.74e- 1
3 EthnicityCaucasian	-12.5	56.7	-0.221	8.26e- 1

If model assumptions are met and `Ethnicity` does not significantly affect average `balance`, there is no need to check

- whether there is significant difference in average `balance` between `Asians` and `African Americans` or between `Caucasians` and `African Americans`

Diagnostics

Diagnostics are the same as those for simple linear regression with a quantitative predictor.



Multiple linear regression

Motivation

- How is `sales` (in thousands of units) for a particular product related to advertising budgets (in thousands of dollars) for TV, `radio` and `newspaper`?
- Model: $\text{sales} = \beta_0 + \beta_1 \times \text{TV} + \beta_2 \times \text{radio} + \beta_3 \times \text{newspaper} + \varepsilon$

We want to examine the relationship between `sales` and budgets for TV, `radio` and `newspaper` *jointly*, instead of *marginally*.

Model

Response Y and p predictors X_1, X_2, \dots, X_p , bound by model

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p + \varepsilon$$

- β_j : change in units in $E(Y)$ for a unit change in X_j while holding all other predictors fixed
- ε : random error term with $E(\varepsilon) = 0$ and $\text{Var}(\varepsilon) = \sigma^2$
- Estimate coefficient vector $\beta = (\beta_0, \beta_1, \dots, \beta_p)$ by the *least squares method*; estimate $\hat{\beta} = (\hat{\beta}_0, \hat{\beta}_1, \dots, \hat{\beta}_p)$ as *LSE (least squares estimate)*

Fitting model

Joint model vs marginal model:

```
# A tibble: 4 x 5
  term      estimate std.error statistic  p.value
<chr>      <dbl>      <dbl>      <dbl>    <dbl>
1 (Intercept)  2.94        0.312        9.42 1.27e-17
2 TV          0.0458      0.00139       32.8 1.51e-81
3 radio       0.189      0.00861       21.9 1.51e-54
4 newspaper -0.00104     0.00587      -0.177 8.60e- 1

# A tibble: 2 x 5
  term      estimate std.error statistic  p.value
<chr>      <dbl>      <dbl>      <dbl>    <dbl>
1 (Intercept) 12.4        0.621       19.9 4.71e-49
2 newspaper   0.0547     0.0166       3.30 1.15e- 3

# A tibble: 2 x 5
  term      estimate std.error statistic  p.value
<chr>      <dbl>      <dbl>      <dbl>    <dbl>
1 (Intercept)  7.03        0.458       15.4 1.41e-35
2 TV          0.0475     0.00269      17.7 1.47e-42
```

Testing on all coefficients

- Is there a relationship between the response and any of the predictors? Namely, is $H_0 : \beta_1 = \beta_2 = \dots = \beta_p = 0$ true?

- Test statistic: F-statistic

$$F = \frac{(TSS - RSS)/p}{RSS/(n - p - 1)},$$

where $TSS = \sum_{i=1}^n (y_i - \bar{y})^2$ with $\bar{y} = n^{-1} \sum_{i=1}^n y_i$ and $RSS = \sum_{i=1}^n (y_i - \hat{y}_i)^2$

- If H_0 is true and the linear model assumptions are correct, F-statistic should be close to 1 on average; under suitable conditions, F-statistic approximately follows an F-distribution

Testing on all coefficients

Testing $H_0 : \beta_1 = \beta_2 = \dots = \beta_p = 0$:

```
value    numdf    dendif
570.2707    3.0000 196.0000
value
1.575227e-96
```

F-statistic: 570.3 with numerator degrees of freedom 3 and denominator degrees of freedom 196; p-value: < 2.2e-16

- Conclusion: reject H_0 , meaning that *at least one of the predictors has a relationship with the response*.

Testing on some coefficients

Is there no relationship between the response and some predictors? Namely, for some $1 \leq q \leq p$, test

$$H_0 : \beta_{p-q+1} = \beta_{p-q+2} = \dots = \beta_p = 0$$

- Fit $M_0 : Y = \beta_0 + \beta_1 X_1 + \dots + \beta_{p-q} X_{p-q} + \varepsilon$, and obtain its residual sum of squares RSS_0
- Fit $M_1 : Y = \beta_0 + \beta_1 X_1 + \dots + \beta_{p-q} X_{p-q} + \dots + \beta_p X_p + \varepsilon$, and obtain its residual sum of squares RSS
- Use test statistic

$$F = \frac{(RSS_0 - RSS)/q}{RSS/(n - p - 1)}$$

Testing on some coefficients

When H_0 and model assumptions are true, test statistic

$$F = \frac{(RSS_0 - RSS)/q}{RSS/(n - p - 1)}$$

approximately follows an F-distribution with numerator degrees of freedom q and denominator degrees of freedom $n - p - 1$

Testing on model fit

- R^2 measures the proportion of variance that is explained by the postulated model
- With three predictors:

```
> FitL3c = lm(sales~TV+radio+newspaper, data=adData)
> summary(FitL3c)$r.squared
[1] 0.8972106
```

- With one predictor:

```
> FitL3d = lm(sales~newspaper,data=adData)
> summary(FitL3d)$r.squared
[1] 0.05212045
```

Interaction terms

Interaction terms: I

Consider predicting the average **sales** (in thousands of dollars) via budgets in advertisement through TV and Radio.

- Model 1: $E(\text{sales}) = \beta_0 + \beta_1 \times \text{TV} + \beta_2 \times \text{Radio}$
- Model 1: how is the change (in unit) in $E(\text{sales})$ relates to a unit change in TV and/or Radio?
- Is model 1 sensible when changes (in unit) in $E(\text{sales})$ are different for a unit change in TV when Radio takes different values?

Interaction terms: I

- If change (in unit) in $E(\text{sales})$ can be different for a unit change in TV at different values of Radio or for a unit change in Radio at different values of TV, then the model

$$E(\text{sales}) = \beta_0 + \beta_1 \times \text{TV} + \beta_2 \times \text{Radio}$$

is no longer suitable

- One way to account for this is to introduce an *interaction* term and use model

$$E(\text{sales}) = \beta_0 + \beta_1 \times \text{TV} + \beta_2 \times \text{Radio} + \beta_3 \times \text{TV} \times \text{Radio}$$

- Does the following model do the job?

$$E(\text{sales}) = \beta_0 + \beta_1 \times \text{TV} + \beta_2 \times \text{Radio} + \beta_3 \times \text{TV}^2 + \beta_4 \times \text{Radio}^2$$

Interaction terms: I

The model

$$E(\text{sales}) = \beta_0 + \beta_1 \times \text{TV} + \beta_2 \times \text{Radio} + \beta_3 \times \text{TV} \times \text{Radio}$$

can be written as

$$E(\text{sales}) = \beta_0 + \beta_1 \times \text{TV} + (\beta_2 + \beta_3 \times \text{TV}) \times \text{Radio}$$

or as

$$E(\text{sales}) = \beta_0 + (\beta_1 + \beta_3 \times \text{Radio}) \times \text{TV} + \beta_2 \times \text{Radio}$$

Interaction terms: I

Fit the model with interaction:

Call:

```
lm(formula = sales ~ TV * radio, data = adData)
```

Residuals:

Min	1Q	Median	3Q	Max
-6.3366	-0.4028	0.1831	0.5948	1.5246

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	6.750e+00	2.479e-01	27.233	<2e-16 ***
TV	1.910e-02	1.504e-03	12.699	<2e-16 ***
radio	2.886e-02	8.905e-03	3.241	0.0014 **
TV:radio	1.086e-03	5.242e-05	20.727	<2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.9435 on 196 degrees of freedom

Multiple R-squared: 0.9678, Adjusted R-squared: 0.9673

F-statistic: 1963 on 3 and 196 DF, p-value: < 2.2e-16


Interaction terms: II

Consider predicting the average **Balance** (of a credit card) using information on if a user is a **Student** (“Yes” or “No”) and his/her **Income**

- Model 0: $E(\text{Balance}) = \beta_0 + \beta_1 \times \text{Income}$
- Model 1: $E(\text{Balance}) = \beta_0 + \beta_1 \times \text{Student} + \beta_2 \times \text{Income}$
- Model 2: $E(\text{Balance}) = \beta_0 + \beta_1 \times \text{Student} + \beta_2 \times \text{Income} + \beta_3 \times \text{Student} \times \text{Income}$

Coding in R: **Student**=“No” is coded as 0 and the baseline, and **Student**=“Yes” as 1

Interaction terms: II



Interaction terms: II

Fit the model with interaction:

Call:

```
lm(formula = Balance ~ Student * Income, data = creditData)
```

Residuals:

Min	1Q	Median	3Q	Max
-773.39	-325.70	-41.13	321.65	814.04

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	200.6232	33.6984	5.953	5.79e-09	***
StudentYes	476.6758	104.3512	4.568	6.59e-06	***
Income	6.2182	0.5921	10.502	< 2e-16	***
StudentYes:Income	-1.9992	1.7313	-1.155	0.249	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 391.6 on 396 degrees of freedom
Multiple R-squared: 0.2799, Adjusted R-squared: 0.2744
F-statistic: 51.3 on 3 and 396 DF, p-value: < 2.2e-16

Diagnostics

Diagnostics

- Diagnostics for multiple linear regression are very similar to those for simple linear regression with a quantitative predictor.
- Additional task: check on collinearity and *variance inflation factor (VIF)*

Collinearity

Collinearity


- refers to the situation in which two or more predictor variables are closely related to each other
- *often inflates the variances of estimated coefficients and makes the model unstable*
- can be measured by the *variance inflation factor (VIF)*

A VIF value that exceeds 5 or 10 indicates a problematic amount of collinearity

Note: $VIF(\hat{\beta}_j) = \frac{1}{1-R_{X_j|X_{-j}}^2}$; collinearity implies $R_{X_j|X_{-j}}^2 \approx 1$

Collinearity

Collinearity among **Limit** and **Rating**:



Collinearity

Model $\text{Balance} \sim \text{Age} + \text{Limit}$:

```
# A tibble: 3 x 5
```

	term	estimate	std.error	statistic	p.value
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	(Intercept)	-173.	43.8	-3.96	9.01e- 5
2	Age	-2.29	0.672	-3.41	7.23e- 4
3	Limit	0.173	0.00503	34.5	1.63e-121

Model $\text{Balance} \sim \text{Rating} + \text{Limit}$:

```
# A tibble: 3 x 5
```

	term	estimate	std.error	statistic	p.value
	<chr>	<dbl>	<dbl>	<dbl>	<dbl>
1	(Intercept)	-378.	45.3	-8.34	1.21e-15
2	Rating	2.20	0.952	2.31	2.13e- 2
3	Limit	0.0245	0.0638	0.384	7.01e- 1

Note: compare standard errors of $\hat{\beta}_{\text{Limit}}$ in both models

Collinearity

```
> FitL3f = lm(Balance~Age+Rating+Limit,data=creditData)
> library(car)
> vif(FitL3f)
```

Age	Rating	Limit
1.011385	160.668301	160.592880

- VIFs indicate considerable collinearity in the data

In case of collinearity, either drop one of the problematic variables or combine some closely related variables

Non-linear models

Non-linear relationship

If there is evidence on a non-linear relationship between response and predictors, we can

- add high-order terms into the model (or employ more advanced non-linear methods); e.g.,

$$E(Y) = \beta_0 + \beta_1 X + \beta_2 X^2$$

- transform predictors (and/or response); e.g., e.g., $E(Y) = \beta_0 + \beta_1 \times f(X)$, where f can be $\log(X)$ or \sqrt{X}

License and session Information

License

```
> sessionInfo()
R version 3.5.0 (2018-04-23)
Platform: x86_64-w64-mingw32/x64 (64-bit)
Running under: Windows 10 x64 (build 19041)

Matrix products: default

locale:
 [1] LC_COLLATE=English_United States.1252
 [2] LC_CTYPE=English_United States.1252
 [3] LC_MONETARY=English_United States.1252
 [4] LC_NUMERIC=C
 [5] LC_TIME=English_United States.1252

attached base packages:
 [1] stats      graphics  grDevices  utils      datasets  methods
 [7] base

other attached packages:
 [1] car_3.0-2      carData_3.0-2 broom_0.5.1   gridExtra_2.3
 [5] ggplot2_3.1.0 knitr_1.21

loaded via a namespace (and not attached):
 [1] tidyselect_0.2.5  xfun_0.4         purrr_0.2.5
 [4] haven_2.0.0       lattice_0.20-35  colorspace_1.3-2
 [7] generics_0.0.2    htmltools_0.3.6  yaml_2.2.0
[10] utf8_1.1.4        rlang_0.4.4      pillar_1.3.1
[13] foreign_0.8-70    glue_1.3.0       withr_2.1.2
```

[16]	readxl_1.2.0	plyr_1.8.4	stringr_1.3.1
[19]	cellranger_1.1.0	munsell_0.5.0	gtable_0.2.0
[22]	zip_1.0.0	evaluate_0.12	labeling_0.3
[25]	rio_0.5.16	forcats_0.3.0	curl_3.2
[28]	fansi_0.4.0	Rcpp_1.0.3	scales_1.0.0
[31]	backports_1.1.3	abind_1.4-5	hms_0.4.2
[34]	digest_0.6.18	openxlsx_4.1.0	stringi_1.2.4
[37]	dplyr_0.8.4	grid_3.5.0	cli_1.0.1
[40]	tools_3.5.0	magrittr_1.5	lazyeval_0.2.1
[43]	tibble_2.1.3	crayon_1.3.4	tidyr_0.8.2
[46]	pkgconfig_2.0.2	data.table_1.11.8	assertthat_0.2.0
[49]	rmarkdown_1.11	rstudioapi_0.8	R6_2.3.0
[52]	nlme_3.1-137	compiler_3.5.0	